

# Deep Reinforcement Learning based on Policy Gradients

---

Ângelo G. Lovatto

`angelo.lovatto@usp.br`

June 27, 2019



**INSTITUTO DE MATEMÁTICA E ESTATÍSTICA**  
UNIVERSIDADE DE SÃO PAULO

# Table of contents

1. Policy Gradients Introduction
2. Vanilla Policy Gradient and GAE
3. Natural Policy Gradient
4. Trust Region Policy Optimization

# Policy Gradients Introduction

---

Environment as a Markov Decision Process (MDP) [7]:  $(\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \rho_0)$ .

- $\mathcal{S}$  is the set of possible states;
- $\mathcal{A}$  is the set of possible actions;
- $\mathcal{P}(s' | s, a)$  is the next state distribution;
- $\mathcal{R}(s, a, s')$  is the reward function;
- $\rho_0(s)$  is the initial state distribution.

# Setup

A *trajectory* is a state-action sequence  $\tau = (s_0, a_0, s_1, \dots, a_{T-1}, s_T)$ , whose probability under a stochastic policy  $\pi$  is given by

$$P(\tau \mid \pi) \doteq \rho_0(s_0) \prod_{t=0}^{T-1} \pi(a_t \mid s_t) \mathcal{P}(s_{t+1} \mid s_t, a_t),$$

and its undiscounted *return* is given by

$$R(\tau) \doteq \sum_{t=0}^{T-1} \mathcal{R}(s_t, a_t, s_{t+1})$$

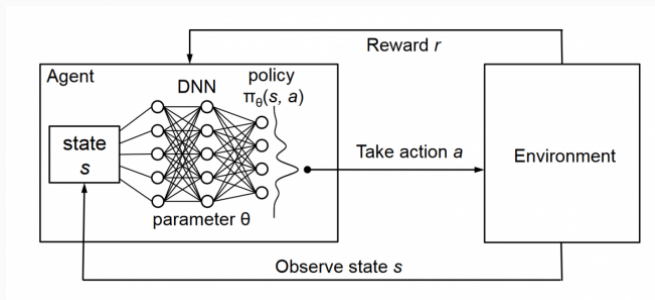
The on-policy value functions are defined as usual:

$$V^\pi(s) \doteq \mathbb{E}_{\tau \sim \pi} [R(\tau) \mid s_0 = s], \quad Q^\pi(s, a) \doteq \mathbb{E}_{\tau \sim \pi} [R(\tau) \mid s_0 = s, a_0 = a].$$

# Policy Gradient Methods - Problem

We consider a smoothly parameterized policy class:

- fix a set of parametric stochastic policies  $\{\pi_{\theta}; \theta \in \mathbb{R}^n\}$
- we assume  $\nabla_{\theta} \pi_{\theta}(a | s), (s, a) \in \mathcal{S} \times \mathcal{A}$ , exists and is finite



We'd like to find the parameters which maximize the expected return

$$\theta^* = \arg \max_{\theta} J(\theta), \text{ where } J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} [R(\tau)].$$

Policy gradient methods optimize the policy using gradient ascent, e.g.,

$$\theta_{k+1} = \theta_k + \alpha \nabla J(\theta_k).$$

Problems:

- how to derive the gradient from the expectation?
- how to obtain sample estimates of the gradient?

# Score Function Estimator

Suppose that  $x$  is a random variable,  $f$  is a function, and we are interested in computing

$$\frac{\partial}{\partial \theta} \mathbb{E}_x [f(x)], \quad x \sim p(\cdot; \theta).$$

The *score function* (SF) estimator of the equation above is derived as follows:

$$\begin{aligned} \frac{\partial}{\partial \theta} \mathbb{E}_x [f(x)] &= \frac{\partial}{\partial \theta} \int p(x; \theta) f(x) dx = \int \frac{\partial}{\partial \theta} p(x; \theta) f(x) dx \\ &= \int p(x; \theta) \frac{\partial}{\partial \theta} \log p(x; \theta) f(x) dx = \mathbb{E}_x \left[ f(x) \frac{\partial}{\partial \theta} \log p(x; \theta) \right]. \end{aligned}$$



# Policy Gradient Derivation

The score function estimator for  $J(\theta)$  is then

$$\begin{aligned}\nabla J(\theta) &= \nabla_{\theta} \mathbb{E}_{\tau \sim \pi_{\theta}} [R(\tau)] \\ &= \mathbb{E}_{\tau \sim \pi_{\theta}} [R(\tau) \nabla_{\theta} \log P(\tau | \theta)],\end{aligned}$$

where

$$\begin{aligned}\nabla_{\theta} \log P(\tau | \theta) &= \nabla_{\theta} \left( \rho_0(s_0) + \sum_{t=0}^{T-1} (\log \pi_{\theta}(a_t | s_t) + \log \mathcal{P}(s_{t+1} | s_t, a_t)) \right) \\ &= \sum_{t=0}^{T-1} \underbrace{\nabla_{\theta} \log \pi_{\theta}(a_t | s_t)}_{\text{no dynamics model needed!}}.\end{aligned}$$

Thus,

$$\nabla J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} R(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]$$

# A Simple Policy Gradient

We can then estimate the gradient by averaging across sampled trajectories:

$$\nabla J(\theta) \approx \hat{g} = \frac{1}{m} \sum_{i=1}^m \sum_{t=0}^{T-1} R(\tau_i) \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}).$$

Intuitively, taking a step in the gradient direction

- increases the probability of trajectories with positive  $R(\tau)$ ,
- decreases the probability of trajectories with negative  $R(\tau)$ .

# Reducing Variance: Temporal Structure

We can remove terms that don't depend on current action:

$$\begin{aligned} & \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} R(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right] \\ &= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \sum_{\substack{t'=0 \\ \text{orange}}}^{T-1} \mathcal{R}(s_{t'+1} | s_{t'}, a_{t'}) \right] \\ &= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \left( \sum_{i=0}^{t-1} \mathcal{R}(s_{i+1} | s_i, a_i) + \sum_{j=t}^{T-1} \mathcal{R}(s_{j+1} | s_j, a_j) \right) \right] \\ &= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \underbrace{\sum_{\substack{t'=t \\ \text{orange}}}^{T-1} \mathcal{R}(s_{t'+1} | s_{t'}, a_{t'})}_{\text{reward-to-go}} \right] \end{aligned}$$

## Reducing Variance: Adding a Baseline

Furthermore, a state-dependent function can be subtracted

$$g \approx \frac{1}{m} \sum_{i=1}^m \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}) \left( \sum_{t'=t}^{T-1} \mathcal{R}(s_{i,t'+1} | s_{i,t'}, a_{i,t'}) - b(s_{i,t}) \right).$$

which is also known as a *control variate*. It can be shown that

$$\mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) b(s_t) \right] = 0,$$

thus, the estimate is still unbiased. This form of is also known as the REINFORCE estimator.

## Vanilla Policy Gradient and GAE

---

In general, the policy gradient can be expressed as

$$g = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} \Phi_t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]$$

where  $\Phi_t$  may be one of the following.

- $R(\tau)$
- $\sum_{k=t}^{T-1} \mathcal{R}(s_{k+1} | s_k, a_k)$
- $\sum_{k=t}^{T-1} \mathcal{R}(s_{k+1} | s_k, a_k) - b(s_t)$
- $Q^{\pi_{\theta}}(s_t, a_t)$
- $A^{\pi_{\theta}}(s_t, a_t)$
- $r_{t+1} + V^{\pi_{\theta}}(s_{t+1}) - V^{\pi_{\theta}}(s_t)$

The advantage function,  $A^{\pi_{\theta}}(s_t, a_t) = Q^{\pi_{\theta}}(s_t, a_t) - V^{\pi_{\theta}}(s_t)$ , yields very low variance in practice.

# Policy Gradient w/ Advantage Function

$$g = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} A^{\pi_{\theta}}(s_t, a_t) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]$$

Intuitively, taking a step in the direction of this gradient:

- increases the probability of better-than-average actions;
- decreases the probability of worse-than-average actions.

We don't have access to the true value functions and sampled rewards have high variance. We can address this in two ways:

- using discounted returns
- using a learned value function (critic)

# Reducing Variance: Discounting

Consider discounted variants of the value functions

$$V^{\pi,\gamma}(s) \doteq \mathbb{E}_{\tau \sim \pi} [R^\gamma(\tau) \mid s_0 = s],$$

$$Q^{\pi,\gamma}(s, a) \doteq \mathbb{E}_{\tau \sim \pi} [R^\gamma(\tau) \mid s_0 = s, a_0 = a],$$

$$A^{\pi,\gamma}(s, a) = Q^{\pi,\gamma}(s, a) - V^{\pi,\gamma}(s),$$

where  $R^\gamma(\tau) \doteq \sum_{t=0}^{T-1} \gamma^t \mathcal{R}(s_t, a_t, s_{t+1})$ .



# Reducing Variance: Discounting

Consider discounted variants of the value functions

$$\begin{aligned}V^{\pi,\gamma}(s) &\doteq \mathbb{E}_{\tau \sim \pi} [R^\gamma(\tau) \mid s_0 = s], \\Q^{\pi,\gamma}(s, a) &\doteq \mathbb{E}_{\tau \sim \pi} [R^\gamma(\tau) \mid s_0 = s, a_0 = a], \\A^{\pi,\gamma}(s, a) &= Q^{\pi,\gamma}(s, a) - V^{\pi,\gamma}(s),\end{aligned}$$

where  $R^\gamma(\tau) \doteq \sum_{t=0}^{T-1} \gamma^t \mathcal{R}(s_t, a_t, s_{t+1})$ .

Note that  $A^{\pi,\gamma}(s, a)$  is equivalent to the *TD residual*,

$$A^{\pi,\gamma}(s, a) = \mathbb{E} [r_{t+1} + \gamma V^{\pi,\gamma}(s_{t+1}) - V^{\pi,\gamma}(s_t) \mid s_t = s, a_t = a].$$

# Reducing Variance: Bootstrapping

In fact, any n-step TD residual is valid.

$$\begin{aligned} A^{\pi, \gamma}(s, a) &= \mathbb{E} [r_{t+1} + \gamma V^{\pi, \gamma}(s_{t+1}) - V^{\pi, \gamma}(s_t) \mid s_t = s, a_t = a] \\ &= \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \gamma^2 V^{\pi, \gamma}(s_{t+2}) - V^{\pi, \gamma}(s_t) \mid s_t = s, a_t = a] \\ &= \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \gamma^3 V^{\pi, \gamma}(s_{t+3}) - V^{\pi, \gamma}(s_t) \mid s_t = s, a_t = a] \\ &\vdots \\ &= \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \cdots + \gamma^{T-1} r_T - V^{\pi, \gamma}(s_t) \mid s_t = s, a_t = a] \end{aligned}$$

# Reducing Variance: Bootstrapping

In practice, we use sample estimates and learned critics

$$\hat{A}_t^{(1)} = r_{t+1} + \gamma V_\phi(s_{t+1}) - V_\phi(s_t)$$

$$\hat{A}_t^{(2)} = r_{t+1} + \gamma r_{t+2} + \gamma^2 V_\phi(s_{t+2}) - V_\phi(s_t)$$

$$\hat{A}_t^{(3)} = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \gamma^3 V_\phi(s_{t+3}) - V_\phi(s_t)$$

$\vdots$

$$\hat{A}_t^{(T)} = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \cdots + \gamma^{T-1} r_T - V_\phi(s_t)$$

Analogous to the n-step return  $G_t^{(n)}$  used in *TD* methods.

# Generalized Advantage Estimation

The generalized advantage estimator is defined as follows

$$\begin{aligned}\hat{A}_t^{\text{GAE}(\gamma, \lambda)} &\doteq (r_{t+1} + \gamma V_\phi(s_{t+1}) - V_\phi(s_t)) && (1 - \lambda) \\ &+ (r_{t+1} + \gamma r_{t+2} + \gamma^2 V_\phi(s_{t+2}) - V_\phi(s_t)) && (1 - \lambda)\lambda \\ &+ (r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \gamma^3 V_\phi(s_{t+3}) - V_\phi(s_t)) && (1 - \lambda)\lambda^2 \\ &\vdots \\ &+ (r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \cdots + \gamma^{T-t-1} r_T - V_\phi(s_t)) && \lambda^{T-t-1}\end{aligned}$$

Analogous to the lambda return  $G_t^\lambda$  used in  $TD(\lambda)$  (forward-view).

It turns out that  $\hat{A}_t^{\text{GAE}(\gamma, \lambda)}$  can be computed as follows

$$\hat{A}_t^{\text{GAE}(\gamma, \lambda)} = \sum_{l=0}^{T-1} (\gamma\lambda)^l \delta_{t+l}^V, \text{ where } \delta_t^V = r_{t+1} + \gamma V(s_{t+1}) - V(s_t).$$

We can recover the one-step TD estimator and the monte carlo by setting  $\lambda$  to 0 and 1 respectively:

$$\begin{aligned}\hat{A}_t^{\text{GAE}(\gamma, 0)} &= \delta_t^V = r_{t+1} + \gamma V(s_{t+1}) - V(s_t), \\ \hat{A}_t^{\text{GAE}(\gamma, 1)} &= \sum_{l=0}^{T-1} \gamma^l \delta_{t+l}^V = R_t^\gamma - V(s_t).\end{aligned}$$

As in  $TD(\lambda)$ ,  $\lambda$  makes a compromise between bias and variance.

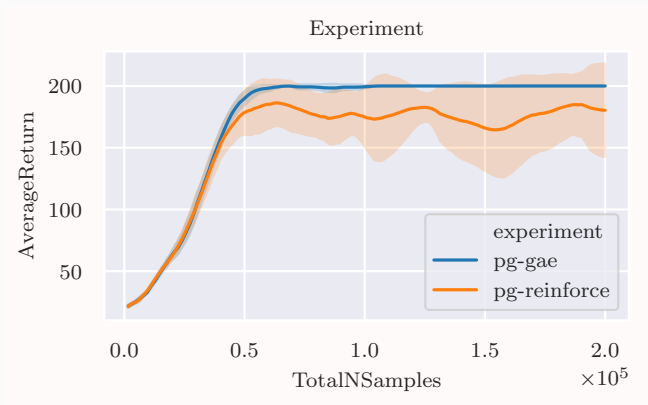
---

**Algorithm 1** Vanilla Policy Gradient + GAE

---

- 1: Initialize policy parameter  $\theta_0$  and value function parameter  $\phi_0$
  - 2: **for**  $k=0, 1, 2, \dots$  **do**
  - 3:     Collect  $N$  transitions  $(s_t, a_t, r_{t+1})$  following  $\pi_{\theta_k}$
  - 4:     Compute  $\delta_t^V$  at all timesteps  $t \in \{1, 2, \dots, N\}$
  - 5:     Compute  $\hat{A}_t = \sum_{l=0}^{T-1} (\gamma\lambda)^l \delta_{t+l}^V$  for every timestep
  - 6:     Fit baseline by minimizing  $\|V_\phi(s_t) - R_t^\gamma\|^2$
  - 7:     Compute  $\theta_{k+1}$  with VPG update
  - 8: **end for**
-

## REINFORCE vs VPG w/ GAE



**Figure 1:** Learning curves in the CartPole-v0 environment, where REINFORCE uses a baseline of zero and VPG uses GAE(0.99,0.97)

# Natural Policy Gradient

---



# Problems with the Vanilla Gradient

$$g = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} A^{\pi_{\theta}}(s_t, a_t) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]$$

- The data distribution depends on the parameters  $\theta$
- Most optimizers work well in settings with i.i.d. data
- Hard to choose step size  $\alpha$ 
  - too big  $\rightarrow$  bad policy  $\rightarrow$  data collected under bad policy  $\rightarrow$  collapse in performance
  - too small: inefficient use of data collected

# Problems with the Vanilla Gradient

Consider a family of policies with parametrization:

$$\pi_{\theta}(a) = \begin{cases} \sigma(\theta) & a = 1 \\ 1 - \sigma(\theta) & a = 2 \end{cases}$$




Figure: Small changes in the policy parameters can unexpectedly lead to **big** changes in the policy.

# Gradient Ascent in Parameter Space

The step size in gradient ascent comes from solving the following optimization problem, e.g., using line search

$$\begin{aligned} d^* &= \arg \max_d J(\theta + d) \\ \text{s.t. } \|d\| &\leq \epsilon \end{aligned}$$

However, the distance in parameters, as we've seen, gives no insight into the resulting policy change.

# Gradient Ascent in Distribution Space

The step size in **natural** gradient ascent comes from solving the following optimization problem,

$$\begin{aligned} d^* &= \arg \max_d J(\theta + d) \\ \text{s.t. } D(\pi_\theta \| \pi_{\theta+d}) &\leq \epsilon \end{aligned}$$

It is better to define the constraint in terms of a "distance" between successive policies.

# Measuring Dissimilarities between Distributions

The Kullback-Leibler (KL) divergence provides a way to measure dissimilarities between distributions.

$$D_{\text{KL}}(P \parallel Q) = \sum_i P(i) \log \frac{P(i)}{Q(i)}$$
$$D_{\text{KL}}(p \parallel q) = \int_{-\infty}^{+\infty} p(x) \log \frac{p(x)}{q(x)} dx$$

We define the KL divergence between policies as follows.

$$\bar{D}_{\text{KL}}(\theta_k \parallel \theta) \doteq \mathbb{E}_{s \sim \pi_{\theta_k}} [D_{\text{KL}}(\pi_{\theta_k}(\cdot \mid s) \parallel \pi_{\theta}(\cdot \mid s))].$$

The objective then becomes

$$\begin{aligned}\theta_{k+1} = \arg \max_{\theta} J(\theta_k) \\ \text{s.t. } \bar{D}_{\text{KL}}(\theta_k \parallel \theta) \leq \delta,\end{aligned}$$

- Easier to pick the distance threshold
- Invariant to the parameterization of the policy
- Hard to calculate in practice

# Natural Policy Gradient: Solution

Using first- and second-order Taylor expansions for the objective and constraint respectively yields

$$\begin{aligned}\theta_{k+1} &= \arg \max_{\theta} J(\theta_k) + g^T(\theta - \theta_k) \\ \text{s.t. } &\frac{1}{2}(\theta - \theta_k)^T F(\theta_k)(\theta - \theta_k) \leq \delta.\end{aligned}$$

where  $F(\theta_k)$  denotes the Hessian of the KL divergence, also known as the Fisher Information Matrix,

$$F(\theta_k) = \nabla_{\theta}^2 \bar{D}_{\text{KL}}(\theta_k \| \theta) |_{\theta=\theta_k}.$$

A Lagrange multiplier argument gives the solution to this constrained optimization problem:

$$\theta_{k+1} = \theta_k + \underbrace{\sqrt{\frac{2\delta}{g^T F(\theta_k)^{-1} g}}}_{\text{scaling term}} \underbrace{F(\theta_k)^{-1} g}_{\text{npg}},$$

# Computing the Natural Gradient

Computing  $F(\theta_k)^{-1}g$  can be expensive:

- DNN policy can have thousands of parameters
- Storing and inverting the Hessian becomes impractical



Computing  $F(\theta_k)^{-1}g$  can be expensive:

- DNN policy can have thousands of parameters
- Storing and inverting the Hessian becomes impractical
- Can instead solve for  $F(\theta_k)x = g$  using conjugate gradient
  - Does not form the whole matrix
  - Requires only a function  $f(v) = F(\theta_k)v$

Computing  $F(\theta_k)^{-1}g$  can be expensive:

- DNN policy can have thousands of parameters
- Storing and inverting the Hessian becomes impractical
- Can instead solve for  $F(\theta_k)x = g$  using conjugate gradient
  - Does not form the whole matrix
  - Requires only a function  $f(v) = F(\theta_k)v$
- We can evaluate  $F(\theta_k)v$  using automatic differentiation.

$$F(\theta_k)v = \nabla_{\theta} \left( (\nabla_{\theta} \bar{D}_{\text{KL}}(\theta_k \| \theta))^{\top} v \right) \Big|_{\theta=\theta_k}.$$

---

## Algorithm 2 Natural Policy Gradient

---

- 1: Initialize policy parameter  $\theta_0$
- 2: **for**  $k=0, 1, 2, \dots$  **do**
- 3:   Collect a set of trajectories following  $\pi_{\theta_k}$
- 4:   Estimate  $\hat{A}_t$ 's using any advantage estimation algorithm
- 5:   Compute the policy gradient  $g$
- 6:   Use CG and Fisher vector products to obtain  $F(\theta_k)^{-1}g$
- 7:   Update the policy parameters

$$\theta_{k+1} = \theta_k + \sqrt{\frac{2\delta}{g^\top F(\theta_k)^{-1}g}} F(\theta_k)^{-1}g,$$

8: **end for**

---

# Trust Region Policy Optimization

---

Natural policy gradient gives us a good step size, however

- How do we ensure improvement in the objective?
- Can we estimate the performance of a new policy?

# Relating Objectives of Two Policies

$$\begin{aligned} J(\tilde{\pi}) - J(\pi) &= J(\tilde{\pi}) - \mathbb{E}_{s_0 \sim \rho_0} [V^\pi(s_0)] \\ &= J(\tilde{\pi}) + \mathbb{E}_{\tau \sim \tilde{\pi}} \left[ \sum_{i=1}^{T-1} V^\pi(s_i) - \sum_{j=0}^{T-1} V^\pi(s_j) \right] \\ &= J(\tilde{\pi}) + \mathbb{E}_{\tau \sim \tilde{\pi}} \left[ \sum_{t=0}^{T-1} V^\pi(s_{t+1}) - V^\pi(s_t) \right] \\ &= \mathbb{E}_{\tau \sim \tilde{\pi}} \left[ \sum_{t=0}^{T-1} \mathcal{R}(s_t, a_t) + V^\pi(s_{t+1}) - V^\pi(s_t) \right] \\ &= \mathbb{E}_{\tau \sim \tilde{\pi}} \left[ \sum_{t=0}^{T-1} A^{\pi}(s_t, a_t) \right] \end{aligned}$$

# Surrogate Objective Function

TRPO makes the following approximation to the result of  $J(\tilde{\pi}) - J(\pi)$ :

$$\mathbb{E}_{\tau \sim \tilde{\pi}} \left[ \sum_{t=0}^{T-1} A^{\pi}(s_t, a_t) \right] \approx \mathbb{E}_{\tau \sim \pi} \left[ \sum_{t=0}^{T-1} \frac{\tilde{\pi}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) \right].$$

That is, we approximate performance of a new policy relative to the old one using data gathered by the latter. In the case of parametric policies, the surrogate objective we seek to optimize is

$$\mathcal{L}(\theta_k, \theta) \doteq \mathbb{E}_{\tau \sim \pi_{\theta_k}} \left[ \sum_{t=0}^{T-1} \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_k}(a_t | s_t)} A^{\pi_{\theta_k}}(s_t, a_t) \right].$$

We could optimize  $\mathcal{L}(\theta_k, \theta)$  freely, however

- we're substituting the trajectory distribution of the old policy for the new policy
- the approximation might not be accurate if the policies differ too much

It turns out that we can bound this approximation error

$$|J(\theta) - \underbrace{(J(\theta_k) + \mathcal{L}(\theta_k, \theta))}_{\approx J(\theta)}| \leq C \sqrt{\bar{D}_{\text{KL}}(\theta_k \| \theta)}$$



TRPO iteratively optimizes the surrogate objective function around a local neighborhood of the current policy, as measured by KL divergence:

$$\begin{aligned} \theta_{k+1} = \arg \max_{\theta} \quad & \mathbb{E}_{\tau \sim \pi_{\theta_k}} \left[ \sum_{t=0}^{T-1} \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_k}(a_t | s_t)} A^{\pi_{\theta_k}}(s_t, a_t) \right] \\ \text{s.t.} \quad & \bar{D}_{\text{KL}}(\theta_k \| \theta) \leq \delta . \end{aligned}$$

The constraint  $\delta$  defines a *trust region* around our local approximation to the new policy's relative performance.

## TRPO: Solution

It can be shown that  $\nabla_{\theta} \mathcal{L}(\theta_k, \theta)$  evaluated at  $\theta = \theta_k$  is equivalent to the policy gradient. Thus, the appropriate Taylor expansions give

$$\begin{aligned} \theta_{k+1} &= \arg \max_{\theta} g^{\top}(\theta - \theta_k) \\ \text{s.t. } &\frac{1}{2}(\theta - \theta_k)^{\top} F(\theta_k)(\theta - \theta_k) \leq \delta, \end{aligned}$$

yielding the same solution as in NPG. TRPO, however, tries to ensure

- improvement in the objective  $J(\theta)$ ,
- satisfaction of the KL constraint.

Final update:

$$\theta_{k+1} = \theta_k + \alpha^j \sqrt{\frac{2\delta}{g^{\top} F(\theta_k)^{-1} g}} F(\theta_k)^{-1} g,$$

where  $\alpha \in (0, 1)$  is a backtracking coefficient.

---

## Algorithm 3 Line Search for TRPO

---

- 1: Compute proposed policy step  $\Delta\theta = \sqrt{\frac{2\delta}{g^\top F(\theta_k)^{-1}g}} F(\theta_k)^{-1}g$
  - 2: **for**  $i = 0, 1, 2, \dots, L$  **do**
  - 3:     Compute proposed update  $\theta = \theta_k + \alpha^i \Delta\theta$
  - 4:     **if**  $\mathcal{L}(\theta_k, \theta) \geq 0$  and  $\bar{D}_{\text{KL}}(\theta_k \parallel \theta) \leq \delta$  **then**
  - 5:         accept the update and set  $\theta_{k+1} = \theta_k + \alpha^i \Delta\theta$
  - 6:         **break**
  - 7:     **end if**
  - 8: **end for**
-

---

## Algorithm 4 TRPO

---



- 1: Initialize policy parameter  $\theta_0$
- 2: **for**  $k=0, 1, 2, \dots$  **do**
- 3:   Collect a set of trajectories following  $\pi_{\theta_k}$
- 4:   Estimate  $\hat{A}_t$ 's using any advantage estimation algorithm
- 5:   Compute the policy gradient  $g$
- 6:   Use CG and Fisher vector products to obtain  $F(\theta_k)^{-1}g$
- 7:   Compute the proposed update  $\Delta\theta = \sqrt{\frac{2\delta}{g^\top F(\theta_k)^{-1}g}} F(\theta_k)^{-1}g$ ,
- 8:   Perform backtracking line search to obtain final update

$$\theta_{k+1} = \theta_k + \alpha^i \Delta\theta$$

9: **end for**

---

# Comparisons



Questions?

# Reward-to-go PG - Proof

Recall that

$$\begin{aligned} & \mathbb{E}_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^{T-1} R(\tau) \nabla_\theta \log \pi_\theta(a_t | s_t) \right] \\ &= \mathbb{E}_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) \sum_{t'=0}^{T-1} \mathcal{R}(s_{t'+1} | s_{t'}, a_{t'}) \right] \\ &= \sum_{t=0}^{T-1} \sum_{t'=0}^{T-1} \mathbb{E}_{\tau \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a_t | s_t) \mathcal{R}(s_{t'+1} | s_{t'}, a_{t'})]. \end{aligned}$$

We'll show that, for  $t' < t$  (when the reward comes before the action),

$$\mathbb{E}_{\tau \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a_t | s_t) \mathcal{R}(s_{t'+1} | s_{t'}, a_{t'})] = 0.$$

## Reward-to-go PG - Proof

Let  $f(t, t') = \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \mathcal{R}(s_{t'+1} | s_{t'}, a_{t'})$ . Note that this term only depends on timesteps  $t, t'$  and  $t' + 1$ . Thus, we can use the *marginal likelihood*

$$\begin{aligned} \mathbb{E}_{\tau \sim \pi_{\theta}} [f(t, t')] &= \int_{\tau} P(\tau | \theta) f(t, t') \\ &= \int_{s_t, a_t, s_{t'}, a_{t'}, s_{t'+1}} P(s_t, a_t, s_{t'}, a_{t'}, s_{t'+1} | \theta) f(t, t') \\ &= \mathbb{E}_{s_t, a_t, s_{t'}, a_{t'}, s_{t'+1} \sim \pi_{\theta}} [f(t, t')] \end{aligned}$$



## Reward-to-go PG - Proof

We can break the expectation using the chain rule of probability:

$$\mathbb{E}_{s_t, a_t, s_{t'}, a_{t'}, s_{t'+1} \sim \pi_\theta} [f(t, t')] = \mathbb{E}_{s_{t'}, a_{t'}, s_{t'+1} \sim \pi_\theta} \left[ \mathbb{E}_{s_t, a_t \sim \pi_\theta} [f(t, t') \mid s_{t'}, a_{t'}, s_{t'+1}] \right].$$

Substituting  $f(t, t')$  in the equation above and noting that the reward term is constant w.r.t. the inner expectation:

$$\begin{aligned} & \mathbb{E}_{s_{t'}, a_{t'}, s_{t'+1} \sim \pi_\theta} \left[ \mathbb{E}_{s_t, a_t \sim \pi_\theta} [f(t, t') \mid s_{t'}, a_{t'}, s_{t'+1}] \right] \\ &= \mathbb{E}_{s_{t'}, a_{t'}, s_{t'+1} \sim \pi_\theta} \left[ \mathbb{E}_{s_t, a_t \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a_t \mid s_t) \mathcal{R}(s_{t'+1} \mid s_{t'}, a_{t'}) \mid s_{t'}, a_{t'}, s_{t'+1}] \right] \\ &= \mathbb{E}_{s_{t'}, a_{t'}, s_{t'+1} \sim \pi_\theta} \left[ \mathcal{R}(s_{t'+1} \mid s_{t'}, a_{t'}) \mathbb{E}_{s_t, a_t \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a_t \mid s_t) \mid s_{t'}, a_{t'}, s_{t'+1}] \right]. \end{aligned}$$

## Reward-to-go PG - Proof

Finally, since  $t' < t$  we can break the inner expectation further using the chain rule of probability

$$\begin{aligned} & \mathbb{E}_{S_t, a_t \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a_t | S_t) | S_{t'}, a_{t'}, S_{t'+1}] \\ &= \mathbb{E}_{S_t \sim \pi_\theta} \left[ \mathbb{E}_{a_t \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a_t | S_t) | S_t] | S_{t'}, a_{t'}, S_{t'+1} \right] \\ &= 0 \end{aligned}$$

Which follows since for any\* probability distribution,

$$\int_x P_\theta(x) \nabla_\theta \log P_\theta(x) = \int_x \nabla_\theta P_\theta(x) = \nabla_\theta \underbrace{\int_x P_\theta(x)}_{=1} = 0.$$

## Second-Order Taylor Expansion of KL

$$D_{\text{KL}}(\theta' \parallel \theta) \approx D_{\text{KL}}(\theta' \parallel \theta') + (\nabla_{\theta} D_{\text{KL}}(\theta' \parallel \theta)|_{\theta=\theta'})^{\top} (\theta - \theta') \\ + \frac{1}{2}(\theta - \theta')^{\top} H(\theta - \theta'),$$

where

$$\nabla_{\theta} D_{\text{KL}}(\theta' \parallel \theta)|_{\theta=\theta'} = - \int_{\mathcal{X}} p_{\theta'}(x) \frac{\nabla_{\theta} p_{\theta}(x)|_{\theta=\theta'}}{p_{\theta'}(x)} \\ = - \int_{\mathcal{X}} p_{\theta'}(x) \nabla_{\theta'} \log p_{\theta'}(x) dx = - \int_{\mathcal{X}} \nabla_{\theta'} p_{\theta'}(x) dx = 0$$

and

$$H = \nabla_{\theta}^2 D_{\text{KL}}(\theta' \parallel \theta)|_{\theta=\theta'} \\ = - \int_{\mathcal{X}} p_{\theta'}(x) \nabla_{\theta'}^2 \log p_{\theta'}(x) dx \\ = - \mathbb{E}_{p_{\theta'}} [\nabla_{\theta}^2 \log p_{\theta'}(x)],$$



J. Achiam, D. Held, A. Tamar, and P. Abbeel.

**Constrained policy optimization.**

In *ICML*, volume 70 of *Proceedings of Machine Learning Research*, pages 22–31. PMLR, 2017.



S. Kakade and J. Langford.

**Approximately optimal approximate reinforcement learning.**

In *ICML*, pages 267–274. Morgan Kaufmann, 2002.



S. M. Kakade.

**A natural policy gradient.**

In *Advances in neural information processing systems*, pages 1531–1538, 2002.



H. Mao, M. Alizadeh, I. Menache, and S. Kandula.

**Resource management with deep reinforcement learning.**

In *Proceedings of the 15th ACM Workshop on Hot Topics in Networks*, pages 50–56. ACM, 2016.



J. Schulman, N. Heess, T. Weber, and P. Abbeel.

**Gradient estimation using stochastic computation graphs.**

In *NIPS*, pages 3528–3536, 2015.



J. Schulman, P. Moritz, S. Levine, M. I. Jordan, and P. Abbeel.

**High-dimensional continuous control using generalized advantage estimation.**

In *ICLR*, 2016.



R. S. Sutton and A. G. Barto.

***Reinforcement learning: An introduction.***

MIT press, 2018.